

ADZUNA: KAGGLE DATA MINING COMPETITION

Braden Simpson
braden@uvic.ca
V00685500

Jordan Ell
jell@uvic.ca
V00660306

March 22, 2013

1 INTRODUCTION

Data mining can be used to solve many complex problems in modern computing, from determining the most. Kaggle¹ is a website that is home to many competitions that solve very complex and intriguing problems. It allows companies such as Facebook, StackOverflow, and others post data science competitions for cash prizes. The competition that is described in this report is job salary prediction, put on by an advertisement search company called Adzuna². By predicting the salary from job posts, Adzuna can provide more accurate ads to their customers.

The competition is about predicting the salary of a job posting based on other attributes in the advertisement. The other attributes are described in detail in Section 2, and with these attributes we will apply algorithms to weigh and train our classifier, from over 240,000 training instances, and then test them on over 40,000 test instances. We use multiple tools and algorithms to obtain a good classifier to this dataset. Section 4 will go in detail about our algorithms and the tools used, as well as the motivation for them. The scores obtained are measured by the classifier's *Mean Absolute Error* (MAE), which is then compared to the rest of the world's submissions to rank them.

The process of our classification is as follows: Data Gathering, Data Preparation, and Data Mining. Gathering consists of collecting the data and interpreting their meaning and schema. Preparation is a large step, as the algorithms

¹A data science competition website. <http://kaggle.com>

²<http://adzuna.com>

and tools require certain specific types of data and schemas, and getting the raw data to this state can be non-trivial, as in the case of a bag-of-words³ method of text processing.

2 DATA GATHERING

The Adzuna competition is all about trying to predict the salaries of jobs posted online from a few choice data points. These predictions will help to bring transparency to the market while helping potential employees discover the best jobs suited for themselves. This being said, the Adzuna team has provided all challengers of the Kaggle competition with data sets in CSV formats.

There are two main data sets used for this competition. The first set contains the training instances. The training instance contains 11 fields of which 10 are attributes with the 11th being the class. The 10 attributes are as follows. An ID provides a unique identifier for each job. This ID will not be used in the actual data mining process but will be used to associate jobs to salaries in the final results of this project. A title is a free text field that gives the position that an employer is looking to fill. A full description is a free text field provided by the job advertisers. This description field holds many of the most interesting potential for future data mining analysis. This field has been given to the competitors stripped of any text that may directly refer to salary. The location raw field is a free text field for the location of the job while the location normalized field is a formal location provided by Adzuna with standard formatting. Contract type provides a categorical look at full time or part time. Contact time provides a categorical view of permanent or contract positions. The company field is a free text field for company names. The job category field is a selection field from 30 standard job categories which employers pick to their needs. Finally, the source name field provides the email address of the employer.

The second set of data provided by the Kaggle competition is the validation set of data. This data again comes in the CSV format and again has 11 fields of which 10 are attributes and 1 as a class. The difference here is that salary has been stripped from the class field. This file is used as a standard testing file for all competitor's data models. This is the file of which all tests will be put against.

There is also an optional file provided by the Kaggle team called the location tree. This file is again a CSV file, however is used for a different purpose. This file is used for the purpose of describing the hierarchical nature of the normalized location field in the training set data. Here the country, direction, city, and region are all given in order to give sense to the previously mentioned normalized location data.

³Bag-of-words is a method of processing unstructured text into a dictionary, without grammar, punctuation, or language nuances

For this project, we have simply made use of both the training and validation sets of data. These pieces of data can be found at the Kaggle website.⁴

3 DATA PREPARATION

Our data mining process intends to use the program Vowpal Wabbit (which will be explained further in the section to follow). Vowpal Wabbit, or VW, takes a custom inout format which needed to be accounted for before any data mining could take place. While preparing this custom data format, various other steps were also taken (to be explained) in order to prepare the data for various attemps and algorithms with data mining. We will outline below the type of input format needed, the custom data transformations, as well as data normalization for future data mining efforts.

The input format for VW is as follows:

$$[Label][Importance[Tag]][NamespaceFeatures|NamespaceFeatures...|NamespaceFeatures] \quad (3.1)$$

where Namespace=String[:Value] and Features=(String[:Value])*. Here we can see the imediate differences to the original CSV files that are provided for the Kaggle challenge. The Label is the floating point number that is being attempted to be predicted. The importance tag can be used to give weight to a specific training instance over the ther instances. This can be useful for weighted preditions or higher confidence in a particular instance. However, due to the large number of training instances given for this Kaggle competition (244,000), the tag was not used. Namespace is a value identifier while features are the values associated with the identifier. In order to transform the CSV files to the given VW input format, python scripts were created and run to make the conversion. While these python scipts were used primarily for conversion purposes, they also allowed the tweaking and adjustments of data as it was converted. We will now explain the different steps taken to tweak data during the conversion.

As it was shown on the Kaggle website, the data given in the training instanced represented a skewed bell curve. In order to get this data into a more acceptable bell curve shape, or normalized (as was needed to VW to run properly, more on this later), two solutions were used, one at a time. The first was simply to take the logarithm of the data set. This creates a more acceptable normalization of the data. That being said, this was the first tweak to be given to the data during the conversion process. The other availible option for creating normalized data is to take the sqaure root of the instances. This also gave a somewhat more normalized view of the data. This tweak was also taken on the data although was preformed at a different time than the logarithm. The results of these two

⁴<http://www.kaggle.com/c/job-salary-prediction/data>

data tweaks can be seen in Table 1 as the different types of normalization transformations.

The second large tweak to be placed on the data during the conversion process was the limiting of free text fields to only keywords. In order to make this tweak, the python library `topia.termextract` was used to extract keywords from local text fields. This means that every text field was analysed sperately for keywords as opposed to looking at all training instances at once. The keyword threshold for reoccurance was set only to 1 as we decided that job descriptions did not always repeat the most important words at any given rate. This data tweak was only applied to data for particular runs of the data mining process and can be seen in the results in Table 1 for how it was used in relation to other data tweaks and their outcomes.

The final data tweak to be placed on the data during the conversion process was the limiting of attributes to be used during the training instance model creation. Our python scripts allowed us to adjust which attributes were to be used for the model generation. Here we played with one major setting in which we turned off all free text fields. The results of turning off the free text fields can again be seen in Table 5.1 in relation with the normalization tweaks. A large number of combinations of free text and categorical attributes could be tested at any given time with or without keywords with this program. However, due to the limited time of this project only the few were actually tested.

This conversion of data to be handles by a larger data set data mining program along with the keyword analysis and other data tweaks represent a non trivial step in our data mining project.

4 DATA MINING

For this project, to reiterate, we are attempting to predict job salaries given some meta information in both categorical and free text forms about the job. The training set provided by Kaggle to issue this challenge is as astounding 244,000 training instances. This being the case, Weka was not an appropriate solution to use as a data mining tool as holding this large amount of data in main memory while preforming standard data mining operatings such as SVM on it would have been devistating. Another issue with this size of data are run times, especially when multiple attempts and different data manipulations will need to be preformed. With these issues present, it is important to either write our own algorithms for speed and size, or use a pre existing tool to handle these scenarios for us. We went with the latter option.

Solving these issues, we found the data mining tool known as Vowpal Wabbit (VW). VW is a command line based data mining tool originally created by Yahoo! Research but was later and more recently sponsored by Microsoft Re-

search. VW was designed to be a fast and lightweight tool from the beginning and accepting of extremely large data sets. VW has a large array of available data mining algorithms, however, its most prominent algorithm (and algorithm used for this project) is known as the sparse gradient descent (GD) on a loss function.

The gradient descent algorithm is a similar structure to the type of gradient descent used in the logistic regression which was taught in the SENG474 data mining class. The main difference here is the use of the weight vector in VW's gradient descent. The weight vector that VW uses has 2^b weights (where b is specified by the b option on the command line). What this really means is that each feature as shown in the data preparation section of the VW input, gets hashed to a particular weight in the vector $[0, 2^b - 1]$. This allows each feature to have its own particular weight in the gradient descent algorithm. There are some complex workings of what happens when two features get hashed to the same weight, but we do not feel the need to get into that in the paper as our feature count was low enough to easily avoid this from happening. (We get roughly 240,000 features in the largest scenarios.) For the loss function of VW, we decided to use the squared loss function. Given a prediction p and a label y a loss function measures the discrepancy between the data model's predicted outcome and the desired outcome. We selected the squared loss function because it has easier input requirements (no real restrictions) than say logistic loss which requires labels of $+1$ or -1 , as well as squared loss allows for a continuous prediction model as opposed to binary predictions. The final mention for the gradient descent method is that input to this method should be normal or close to. The problem with unnormalized data comes when a step is made in the gradient descent. At first the (hypothetical) slope of the gradient is so high that large steps cause an overshooting of the best target value, or the well of the parabola. Smaller first steps followed by larger later steps can help the algorithm hone in on the well of the parabola. Thus having a bell shape of the input data, or having it normalized, gives us a better process towards finding the well of the parabola.

VW comes with its own unique input format as was described in the data preparation section of this paper. However what was not fully explained was the flexibility of the input and how the input's features are interpreted in the data mining process. (Refer to Equation 1) for the following explanation.) The label is the most important feature of the input, which is the class attribute of the training instance. Since VW is expecting a normalized input in terms of the distribution of class values in the training instance, it is important for this tool that the training instance be normalized with either the logarithm function or the square root function before running.

For the actual data mining process, 6 main data trials were given as input to the VW program. These trials and their data varied as follows. First, a simplistic approach was taken by excluding all bag of words columns in the

training instances. Bag of words form the most complex data mining rules and therefore are excluded to create a base line analysis of where the tool will finish without any of the added complexity. This data set trial was used with both the logarithmic first and square root functions of data normalization. (All data trials from now on are used once with logarithmic and once with square root normalization.) These results of these base line trials runs can be found in Table 5.1 in the last section of the paper. The next trial to be run on the data mining algorithm (gradient descent) was what we call the full bag of words run. Here, all columns that are bag of words and not categorical were used in addition to all categorical columns. The main note here is that the bag of words columns were pre processed in order to eliminate numbers, abbreviations, or any other text that cannot be classified in a standard dictionary. This pre processing only leaves English words in the free text columns of the training instances. The results of this trial run can again be found in Table ???. The final, and most complex trial run against the gradient descent algorithm was what we are calling bag of keywords. As the title of this trial gives away, we were only concerned with keywords in the free text columns of the training instances. In order to get the keywords of a particular bag of text, the python library known as Topia⁵ was used. Topia is a natural language processor library which can (with some flexibility) extract the keywords from a block of text. An example of keywords are as follows. The text "The fox can't jump over the fox's tail." in Topia yields the results: [('tail', 1, 1), ('fox', 2, 1)]. As it can easily be seen, the keywords of tail and fox have been extracted from the main body of text as keywords. What is less obviously are the remaining digits provided by Topia. This is where the flexibility of Topia comes into play. In each triple, the keyword, the occurrences, and the number of words in the keyword are given. This being the case, we are able to set thresholds on Topia's extractions in order to limit the keywords selected. We can limit the keywords to only those repeated in the text body at least once, or those that are composed of at least two words and so on. However, for simplicity of our main data trials, we allowed topia the lowest of thresholds for keyword extractiong, being no repeat and single words. This allowed us to again create a base line measurement of keyword usage in the data mining process. Again pre processing was done on the training instances in order to limit the free text columns to key words only. The results of this baseline keyword run in VW can be seen in Table ??.

5 RESULTS AND CONCLUSIONS

This section contains the main results table for the 6 main data trial runs as described in the data mining section of this paper. All results are expressed in mean average error (MAE). The MAE is calculated when uploading potential results to the Kaggle website competition. The MAE is exactly what it sounds like, in that it is the average "dollars off" between our estimate created using

⁵<https://pypi.python.org/pypi/topia.termextract/>

our data mining model and the actual answer known only by Kaggle.

Free Text	Logarithmic	Square Root
None	12712.17	13293.12
All	7350.44	7175.04
Keywords		

Table 5.1: Results from main data trial runs.

As can be seen in the table above, the baseline scenario of no bag of words used in the data mining process yielded by far the worst results. This was largely expected based on the nature of the Kaggle competition. If data mining was better without bag of word usage, these types of challenges may not be ever needed as the simplisity of data mining would be trivial. Where the power and knowledge of this competition really comes from is the use of the bag of words or free text columns in the training instances and how they can be used to better the results of any data mining predictions.

The next man data trial run was that on the full bag of words or free text columns inside the VW data mining tool. These results are vastly superior to those of the non free text trials. This was also largely to be expected. Some notes to take away from these results are as follows. For one, the use of all words in free text columns and the VW data mining tool using the gradient descent algorithm scored a lower MAE than the benchmark set by Kaggle which was a random forest algorithm. The Kaggle benchmark with the random forest was 7536.29 which was beat by the full bag of words run we preformed. This already is showing promise for the use of the VW tool over a standard tool such as Weka which may just be able to load the data and preform analysis. The Kaggle random forest benchmark was also said to take 1.5 hours on a 2.7GHz processor the 8 cores and 8GB of RAM. The VW tool took around 5 minutes to fully complete its run cycle. The last thing to note from this full bag of words trial is that it put our team into 108th spot on the Kaggle leaderboards for this challenge out of 253 participating teams. This is a large success for this project.

The final result for the main data trial sets was that of BLANK and BLANK from the keyword selection of the bag of words columns in the training instances. Write something here when data is done.

In conclusion, write something here once all data is done.